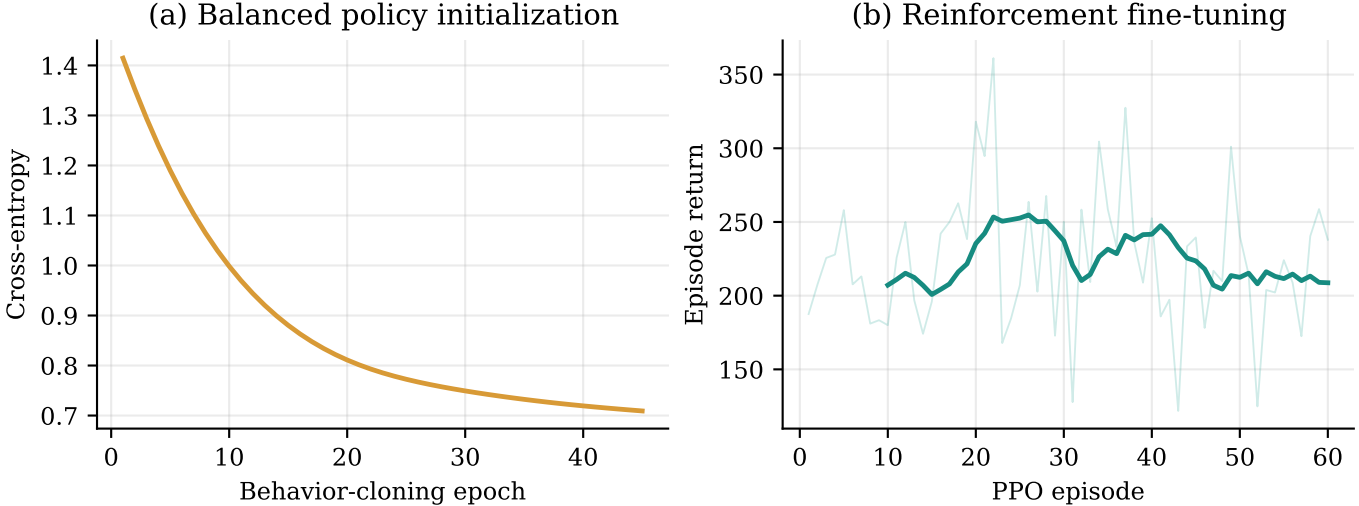
(a) Balanced policy initialization
Cross-entropy
1.4
1.3
1.2
1.1
1.0
0.9
0.8
0.7
0
10
20
30
40
Behavior-cloning epoch
(b) Reinforcement fine-tuning
Episode return
350
300
250
200
150
0
10
20
30
40
50
60
PPO episode